



CS61C

Great Ideas
in
Computer Architecture
(a.k.a. Machine Structures)



Teaching Professor
Dan Garcia

C Pointers, Arrays, and Strings



Agenda

Pointers

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings




Memory is a Single Huge Array

- Consider memory to be a byte-addressed array.
 - Each cell of the array has an **address** associated with it.
 - Each cell also stores some **value**.
 - For now, with this abstraction, let's us think we have access to ∞ memory, numbered from 0...

Addresses are commonly written in hexadecimal format.

Location (address)

... 0x100 0x101 0x102 0x103 0x104 0x105 0x106 0x107 0x108 0x109 0x10A 0x10B ...



C variable names

x


c

p



Pointers store addresses

- Consider memory to be a byte-addressed array.
 - Each cell of the array has an **address** associated with it.
 - Each cell also stores some **value**.
 - For now, with this abstraction, let's think we have access to ∞ memory, numbered from 0...
- Pointer**: A variable that contains the address of another variable.
 - In other words, it "**points**" to a memory location.





Pointer Syntax

0x100

p

???

0x104

x

3

```
1 int *p;  
2 int x = 3;  
  
3 p = &x;  
  
4 printf("p points to %d\n",  
        *p);  
  
5 *p = 5;
```

- Declaration
- Tells compiler that **variable p is address** of an int

Pointer Syntax



```
1 int *p;  
2 int x = 3;  
3 p = &x;  
4 printf("p points to %d\n",  
        *p);  
5 *p = 5;
```

- Declaration
- Tells compiler that variable `p` is address of an `int`
- Tells compiler to assign **address of `x`** to `p`
- `&`: "**address operator**" in this context

Pointer Syntax



1 `int *p;`

2 `int x = 3;`

3 `p = &x;`

4 `printf("p points to %d\n",

*p);`

5 `*p = 5;`

- Declaration
- Tells compiler that variable `p` is address of an `int`
- Tells compiler to assign **address of `x`** to `p`
- `&`: "**address operator**" in this context
- Gets **value pointed to** by `p`
- `*`: "**dereference operator**" in this context

Pointer Syntax



1 `int *p;`

2 `int x = 3;`

3 `p = &x;`

4 `printf("p points to %d\n",
 *p);`

5 `*p = 5;`

- Declaration
- Tells compiler that variable `p` is address of an `int`
- Tells compiler to assign **address of `x`** to `p`
- `&`: "**address operator**" in this context
- Gets value pointed to by `p`
- `*`: "**dereference operator**" in this context
- Changes **value pointed to** by `p`
- Use deref operator `*` on left of `=`



Pointer Syntax



The "*" is used in two ways:

Declaration (L1): Indicate p is a pointer

Dereference (L4,L5): Value pointed to by p

```
1 int *p;
2 int x = 3;

3 p = &x;

4 printf("p points to %d\n",
        *p);

5 *p = 5;
```

- **Declaration**
- Tells compiler that variable p is address of an int
- Tells compiler to assign **address of x** to p
- &: "**address operator**" in this context
- Gets **value pointed to** by p
- *: "**dereference operator**" in this context
- Changes **value pointed to** by p
- Use deref operator * on left of =



Pointers are Useful When Passing Parameters

C is **pass-by-value**: A function parameter gets assigned **a copy of the argument value**.

Changing the function's copy cannot change the original.

```
void addOne (int x)
{
    x = x + 1;
}
```

x 3 4

y 3

```
int y = 3;
addOne(y);
```

▲ y is still 3...



Pointers are Useful When Passing Parameters

C is **pass-by-value**: A function parameter gets assigned **a copy of the argument value**.

Changing the function's copy cannot change the original.

```
void addOne (int x)
{
    x = x + 1;
}
```

```
int y = 3;
addOne(y);
```

x 3 4

y 3

⚠ y is still 3...

To get a function to change a value, **pass in a pointer**.

```
void addOne (int *p)
{
    *p = *p + 1;
}
```

```
int y = 3;
addOne(&y);
```

p 0x100

0x100
y 3 4

✓ y is now 4!



Pointers in C ... The Good, the Bad, and the Ugly

- To pass a large struct or array to a function, it's **easier/faster/etc.** to pass a pointer.
 - Otherwise, we'd need to copy a huge amount of data!
- At the time C was invented (early 1970s), compilers didn't produce efficient code, so C was designed to give human programmer more flexibility.
 - Nowadays, computers are 100,000x faster; compilers are also way, way, way better.
- Still used for low-level system code, as well as implementation of "pass-by-reference" object paradigms in other languages.
- In general, pointers allow cleaner, more compact code.



Pointers in C ... The Good, the Bad, and the Ugly

⚠ So, what are the drawbacks?

- Pointers are probably the single largest source of bugs in C.


Be careful!

- Most problematic with dynamic memory management
 - Dangling references and memory leaks
- } (more later)

Common C Bug: Garbage Addresses

- Declaring a pointer just allocates space to hold the pointer.
 - It does not allocate something to be pointed to!
- Recall: Local variables in C **are not initialized**.
 - They may contain anything.
- What does the following code do?

```
void f()
{
    int *ptr;
    *ptr = 5;
}
```





Agenda

Using Pointers Effectively

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings

Pointers to Different Data Types

- Pointers are used to point to a variable of a particular data type.
 - Normally a pointer can only point to one type.

```
int *xptr;  
char *str;  
struct llist *foo_ptr;
```

(more later, but)

- **void *** is a type that can point to anything (**generic pointer**).
 - Use sparingly to help avoid program bugs... and security issues...and more...
- You can even have pointers to functions...
 - `int (*fn) (void *, void *) = &foo;`
 - `fn` is a function that accepts two `void *` pointers and returns an `int` and is initially pointing to the function `foo`.
 - `(*fn) (x, y);` will then call the function



NULL pointers...

The pointer of all 0s is special.

- The **NULL** pointer, like in Java, Python, etc...

```
char *p = NULL;  
p 0x00000000
```

If you write to or read from a null pointer, your program should **crash**.

Since **0 is false**, its very easy to do tests for null:

- `if(!p) { /* p is a null pointer */ }`
- `if(q) { /* q is not a null pointer */ }`

More Typing: Typedefs and Structs

- typedef allows you to define new types.

```
typedef uint8_t BYTE;
BYTE b1, b2;
```

- structs are structured groups of variables, e.g.,

```
typedef struct {
    int length_in_seconds;
    int year_recorded;
} SONG;
```

} SONG is an alias for typedef struct {int length_in_seconds; int year_recorded; }

```
SONG song1;
song1.length_in_seconds = 213;
song1.year_recorded = 1994;
```

Dot notation: x.y = value

song1

length_in _seconds	213	year_ recorded	1994
-----------------------	-----	-------------------	------

```
SONG song2;
song2.length_in_seconds = 248;
song2.year_recorded = 1988;
```

Structs are not objects!
The **dot (.)** operator is **not** a method call! (more later)

Struct Pointers

```
typedef struct {  
    int x;  
    int y;  
} Coord;
```

```
/* declarations */  
Coord coord1, coord2;  
Coord *ptr1, *ptr2;
```


```
/* instantiations  
go here... */
```

...

```
/* dot notation */  
int h = coord1.x;  
coord2.y = coord1.y;
```

```
/* arrow notation = deref + struct access*/  
int k;  
k = (*ptr1).x;  
k = ptr1->x;  // equivalent
```

```
/* This compiles, but what does it do? */  
ptr1 = ptr2;
```




Garcia, FA25


Assume ptr2 points to an initialized Coord struct {3, 4}. What does ptr1 = ptr2; do?



CHOICE A
(click on this side)




CHOICE B
(click on this side)






CHOICE A
(click on this side)



CHOICE B
(click on this side)






Struct Pointers

```
typedef struct {  
    int x;  
    int y;  
} Coord;  
  
/* declarations */  
Coord coord1,  
coord2;  
Coord *ptr1, *ptr2;  
  
/* instantiations  
go here... */
```

```
/* dot notation */  
int h = coord1.x;  
coord2.y = coord1.y;  
  
/* arrow notation = deref + struct access*/  
int k;  
k = (*ptr1).x;  
k = ptr1->x; // equivalent  
  
/* This compiles, but what does it do? */  
ptr1 = ptr2;
```



Garcia, FA25



Agenda

Arrays, Pointer Arithmetic

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings



A C array is really just a big block of memory

Declaration:

- `int arr[2];`
- ...declares a 2-element integer array



arr

Declaration and initialization

- `int arr[] = {795, 635};`
- declares and fills a 2-elt integer array



arr

Accessing elements:

- `arr[num]`
- returns the numth element.
- This is shorthand for **pointer arithmetic**.

`arr[0];` // 795



Pointer Arithmetic

Equivalent:


$a[i] \Leftrightarrow *(a+i)$

pointer + n

- Adds $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$ to memory address.

pointer - n

- Subtracts $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$ from memory address.



```
// 32-bit unsigned int array  
uint32_t arr[] = {50, 60, 70};
```

```
uint32_t *q = arr;
```

Pointer Arithmetic

Equivalent:

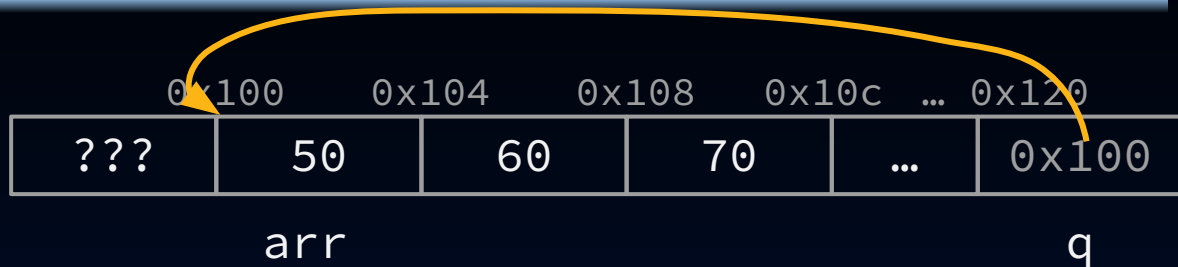
$a[i] \Leftrightarrow *(a+i)$

pointer + n

- Adds $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$ to memory address.

pointer - n

- Subtracts $n * \text{sizeof}(\text{"whatever pointer is pointing to"})$ from memory address.



```
// 32-bit unsigned int array
uint32_t arr[] = {50, 60, 70};
```

```
uint32_t *q = arr;
```

```
printf("    *q: %d is %d\n", *q, q[0]);
printf("*(q+1): %d is %d\n", *(q+1), q[1]);
printf("*(q-1): %d is %d\n", *(q-1), q[-1]);
```

```
*q: 50 is 50
*(q+1): 60 is 60
*(q-1): ??? is ???
```



How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

```
1 void increment_ptr(int32_t *p)
2 {   p = p + 1;   }

3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(q);
6 printf("*q is %d\n", *q);
```

Handling Pointers (1/2)

How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is **pass-by-value**!

```


1 void increment_ptr(int32_t *p)
2 { p = p + 1; }

3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(q);
6 printf("*q is %d\n", *q);

```

0x100

p



arr

q

Handling Pointers (1/2)

How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is **pass-by-value**!

```

1 void increment_ptr(int32_t *p)
2 { p = p + 1; }


3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(q);
6 printf("*q is %d\n", *q);
    
```

0x100
0x104

p

Print output:

***q is 50**



arr

q



Handling Pointers (2/2)

How to get a function to change a pointer?

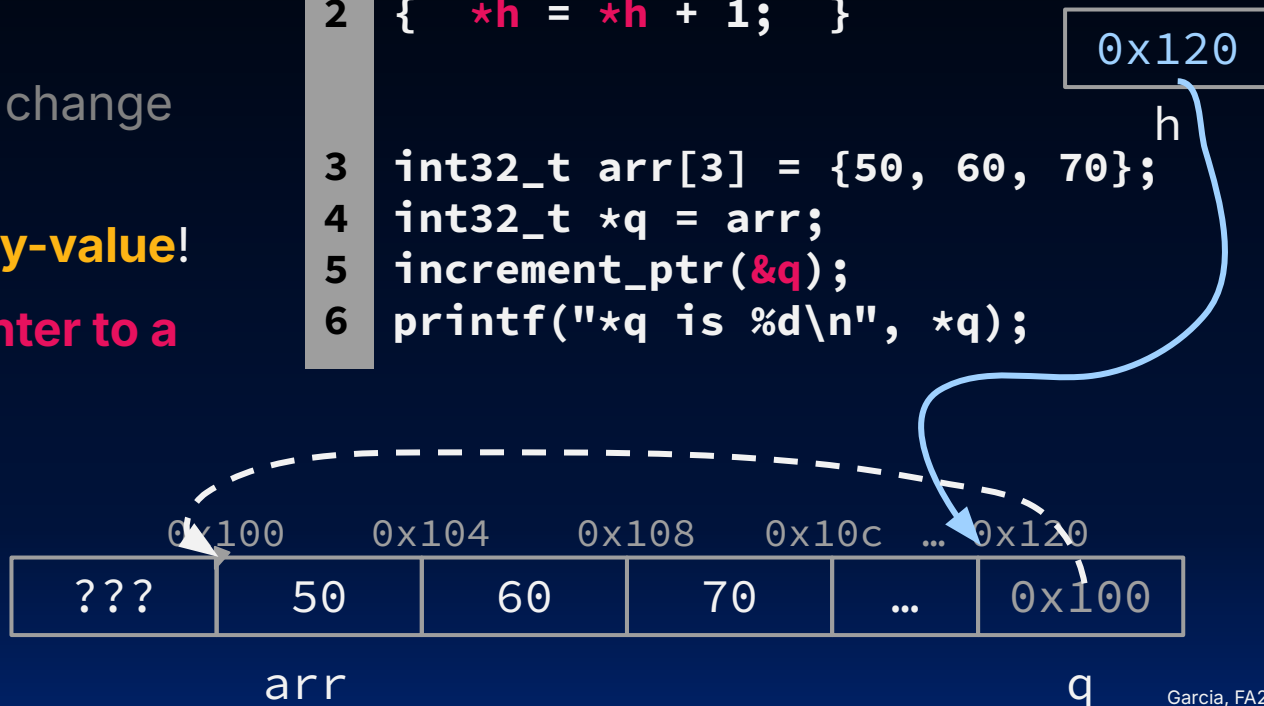
- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is **pass-by-value**!

- Instead, pass **a pointer to a pointer** ("handle").
- Declared as `data_type **h`.

```
1 void increment_ptr(int32_t **h)
2 { *h = *h + 1; }

3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(&q);
6 printf("*q is %d\n", *q);
```





Handling Pointers (2/2)

How to get a function to change a pointer?

- Suppose we want `increment_ptr` to change where `q` points to.

Remember: C is **pass-by-value**!


- Instead, pass **a pointer to a pointer** ("handle").
- Declared as `data_type **h`.

Print output:

`*q is 60`

```
1 void increment_ptr(int32_t **h)
2 { *h = *h + 1; }

3 int32_t arr[3] = {50, 60, 70};
4 int32_t *q = arr;
5 increment_ptr(&q);
6 printf("*q is %d\n", *q);
```



arr

q



Agenda

Array Pitfalls

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings

Array Pitfall #1 of 123234983

Declare array and initialize all elements of an array of known size **n**:

- 

```
int i, arr[10];
for(i = 0; i < 10; i++){ ... }
```

- Strongly encouraged

```
int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```

Why? SINGLE SOURCE OF TRUTH!

- Utilize **indirection** and avoid maintaining two copies of the number 10!

Arrays vs Pointers

Arrays are (almost) identical to pointers.

- `char *string` and `char string[]` are nearly identical declarations
- They differ in **very subtle** ways: incrementing, declaration of filled arrays...(more in a bit)


Accessing Array Elements

- `arr` is an array, but it looks like a pointer in many respects (though not all).
- `arr[0]` is the same as `*arr`
- `arr[2]` is the same as `*(arr+2)`



Arrays are not implemented as you'd think...


```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```





Arrays are not implemented as you'd think...


```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```



```
*p = 1;  
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d: signed decimal, %x: hex
```

```
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```





Arrays are not implemented as you'd think...


```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```

```
*p:1, p:108, &p:100
```

```
*p = 1;  
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d: signed decimal, %x: hex
```

```
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```





Arrays are not implemented as you'd think...


```
int *p, *q, x;  
int a[4];  
p = &x;  
q = a + 1;
```

```
*p:1, p:108, &p:100  
*q:2, q:110, &q:104
```

```
*p = 1;  
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d: signed decimal, %x: hex
```

```
*q = 2;  
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;  
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```



Arrays are not implemented as you'd think...

```
int *p, *q, x;
int a[4];
p = &x;
q = a + 1;
```

```
*p = 1;
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d: signed decimal, %x: hex
```

```
*q = 2;
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```


```
*p:1, p:108, &p:100
*q:2, q:110, &q:104
*a:3, a:10c, &a:10c
```



K&R:

"An array name is not a variable"

Us: "A C array is really just a big block of memory"





Agenda

Strings

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings



C Strings

[READ ON YOUR OWN]


A C string is just an array of characters, followed by **a null terminator**.

- Null terminator: the byte 0 (number) aka `'\0'` character)

The standard C library `string.h` assumes **null-terminated strings**.

- String operations often abstract away null terminators when returning values!

```
char str[] = "abc";
```



```
... strlen(str) ... // 3 ⚠
```

```
// possible implementation*  
int strlen(char s[])  
{  
    int n = 0;  
    while (*(s++) != 0) { n++; }  
    return n;  
}
```

*for actual `strlen()` implementation, see glibc

Lab 01: `string.h`: `strcpy()` vs `strncpy()`



And in Conclusion...

- C pointers and arrays are **pretty much the same**, except with function calls.
- C knows how to **increment pointers**.
- C is an efficient language, but with little protection.
 - Array bounds **not checked**
 - Variables **not automatically initialized**
- Use handles to change pointers
- Strings are arrays of characters with a null terminator
 - The length is the # of characters, but memory needs 1 more for `\0`
- (Beware) The cost of efficiency is more overhead for the programmer.
 - "C gives you a lot of extra rope, don't hang yourself with it!"



More C Features

Full slides for reference in
[L03](#).

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings



From Last Time

[from last time]

```
1 #include <stdio.h>
2 int main(int argc, char *argv[]) {
3     int x = 0;
4     int y;
5
6     printf("before: x=%d, y=%d\n",
7           x, y);
8     x++;
9     y += x; // y = y + x
10    printf(" after: x=%d, y=%d\n",
11           x, y);
12    return 0;
13 }
```

%d Placeholder for argument, where
d: format value as decimal numeral

Print output:

```
before: x=0, y=22621
after:  x=1, y=22622
```

Declaring a C Variable Does Not Also Initialize It

- Variables are **not automatically initialized to default values!**
- If a variable is not initialized in its declaration, **it holds garbage!**
 - The contents are undefined...
 - ...but C still lets you use uninitialized variables!
- ⚠ Danger ⚠**: Bugs sometimes may only manifest after you've built other parts of your program.

```
int x; // declaration
...
x = 42; // initialization
```

x ???? 42

```
// OK
int y = 38; y
```

38




Why are variables typed in C?

[from last time]

- C variables are typed.
 - **Types of variables** can't change, e.g., `y` cannot store a float.
 - However, you can **typecast values** (more later)
- A variable's type helps the compiler determine **how to translate the program to machine code** designed for the computer's architecture:
 - **How many bytes** the variable takes up in memory, and
 - **What operators** the variable supports, etc.

```
uint16_t y = 38;
```

(unsigned 16-bit integer,
see `stdint.h`)





The C bool: true or false?

[from last time]

- Originally, the boolean was not a built-in type in C! Instead:
- FALSE:
 - **0** (integer, i.e., all bits are 0)
 - **NULL** (pointer) (more later)
- TRUE:
 - **Everything else!**
 - (Note: Same is true in Python)
- Nowadays:
 - **true** and **false** provided by `stdbool.h`.
 - Built-in type as of C23 (but we are using C17)

```
if (42) {  
    printf("meaning of life\n");  
}
```

[THESE WERE MOVED TO THE END <END>]

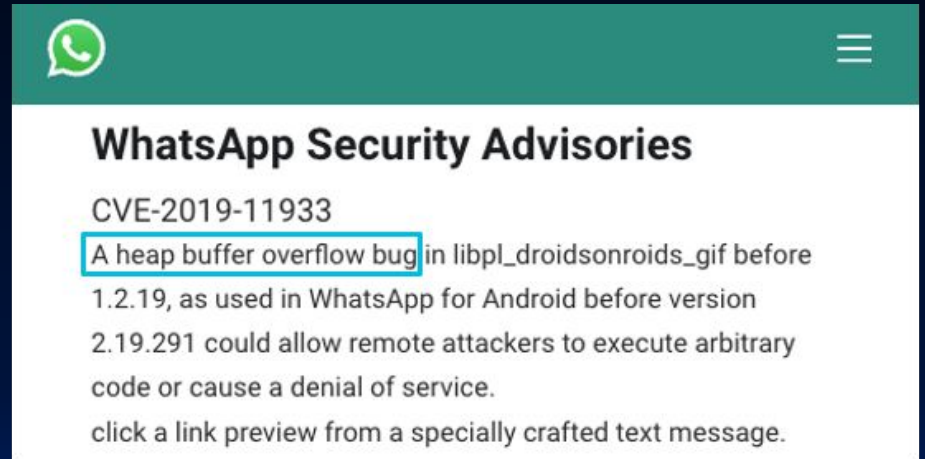


Arrays Are Very Primitive (1/3)

[GOT TO HERE]

1. Array bounds are not checked during element access.
 - Consequence: We can accidentally access off the end of an array
 - Corrupts other parts of program...
 - Including internal C data
 - May cause crashes later.

```
int N = 100;
int foo[N];
int i;
....
for(i = 0; i <= N; ++i) {
    foo[i] = 0;
}
```



"Buffer overflow"

Arrays Are Very Primitive (2/3)

[for next time]

2. An array is passed to a function as a pointer.
 - o Consequence: The array size is lost! Be careful with sizeof()!

same as `int *arr`

```
int bar(int arr[], unsigned int size )
{
    ... arr[size - 1] ...
}

int main(void)
{
    int a[5], b[10];
    ...
    bar(a, 5);
    ...
}
```

You should always explicitly include array length as a parameter.





Arrays Are Very Primitive (3/3)

[for next time]

3. Declared arrays are only allocated while the scope is valid.

```
// incorrect
char *foo() {
    char string[32]; ...;
    return string;
}
```

Solution:
Dynamic memory
allocation!

(more later)



Arrays Are Very Primitive, Summary

[for next time]

1. Array bounds are not checked during element access.
2. An array is passed to a function as a pointer.
3. Declared arrays are only allocated while the scope is valid.

Consequences? Segmentation faults, bus errors, ...

- These are VERY difficult to find; be careful!
- You'll learn how to debug these in Lab 02 with gdb...



Strings

- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings



Memory, Words, and Endianness

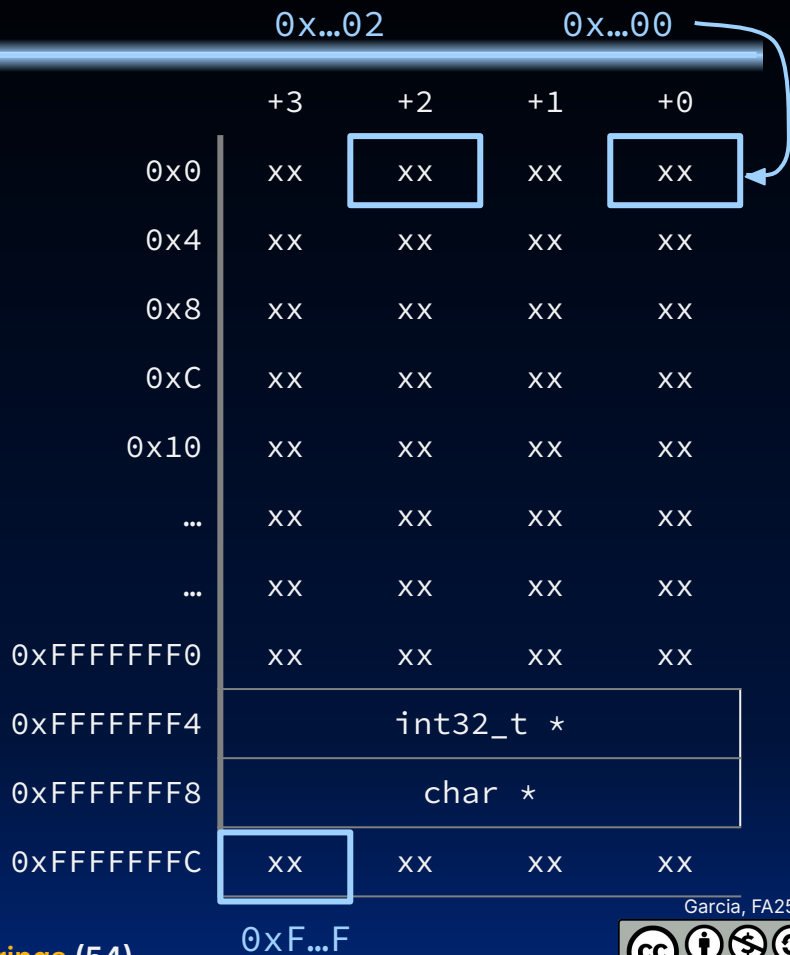
- Pointers
- Using Pointers Effectively
- Arrays, Pointer Arithmetic
- Array Pitfalls
- Strings



Memory and Addresses

How to read byte addresses:

- Modern machines are "byte-addressable."
 - Hardware's memory composed of 8-bit storage cells; each byte has a unique address
- We commonly think in terms of **word size**:
 - aka **number of bits in an address**
- A 32b architecture has:
 - 4B words
 - word-sized pointers = 4B pointers
 - `sizeof(int *) == ...`
 - `== sizeof(char *) == 4`



García, FA25



Example program memory, on hive

see: endianness.c

```
#include <stdio.h>
#include <stdint.h>

int main(int argc, char *argv[]) {
    int32_t value = 0x12345678;
    char str[] = "bye";
    char str2[] = "hello";
    int16_t short_val = 0xaabb;
    ...

    return 0;
}
```

	+3	+2	+1	+0
0x0	xx	xx	xx	xx
0x4	xx	xx	xx	xx
...	xx	xx	xx	xx
0x7F...FE164	0xaa	0xbb	xx	xx
0x7F...FE168	0x12	0x34	0x56	0x78
0x7F...FE16C	'y'	'b'	xx	xx
0x7F...FE170	'e'	'h'	'\0'	'e'
0x7F...FE174	'\0'	'o'	'l'	'l'
0x7F...FE178	xx	xx	xx	xx
...	xx	xx	xx	xx
0xFFFFFFFFFC	xx	xx	xx	xx

Garcia, FA25



Word Alignment

- We often want **word alignment**:
 - Some 32-bit processors will not allow you to address 32b values without being on 4-byte boundaries.
 - Others will just be very slow if you try to access "unaligned" memory
- Based on this diagram, does the hive machine processor do word-alignment?

	+3	+2	+1	+0
0x0	xx	xx	xx	xx
0x4	xx	xx	xx	xx
...	xx	xx	xx	xx
0x7F...FE164	0xaa	0xbb	xx	xx
0x7F...FE168	0x12	0x34	0x56	0x78
0x7F...FE16C	'y'	'b'	xx	xx
0x7F...FE170	'e'	'h'	'\0'	'e'
0x7F...FE174	'\0'	'o'	'l'	'l'
0x7F...FE178	xx	xx	xx	xx
...	xx	xx	xx	xx
0xFFFFFFFFC	xx	xx	xx	xx

Garcia, FA25



Endianness

- The hive machines are **little endian**.

- Words are stored **least significant byte (LSB) first**.

- (contrast: big endian, most significant byte (MSB) stored first)

- LSB of integers have lowest byte address

`int32_t value = 0x12345678;`

- Diagram makes it easy to read integers, harder to read strings:

`char str[] = "bye";`

<u>'b'</u>	'y'	'e'	'\0'
------------	-----	-----	------

	+3	+2	+1	+0
0x0	xx	xx	xx	xx
0x4	xx	xx	xx	xx
...	xx	xx	xx	xx
0x7F...FE164	0xaa	0xbb	xx	xx
0x7F...FE168	0x12	0x34	0x56	0x78
0x7F...FE16C	'y'	'b'	xx	xx
0x7F...FE170	'e'	'h'	'\0'	'e'
0x7F...FE174	'\0'	'o'	'l'	'l'
0x7F...FE178	xx	xx	xx	xx
...	xx	xx	xx	xx
0xFFFFFFF0	xx	xx	xx	xx

Garcia, FA25



For you to try at home with gdb

[at home]
endianness.c

```
$ gdb endianness
(gdb) b 12 # set breakpoint at line 12
(gdb) r    # run, initializing all vars
(gdb) p/x str # see string bytes
(gdb) p/x str2
(gdb) # print one word, show in hex
(gdb) x/1wx 0x7fffffffef164
(gdb) <enter> # repeat last action
(gdb) ... # keep pressing <enter>
(gdb) # LSB in lowest address/word
(gdb) # print 4 bytes, show in hex
(gdb) x/4bx 0x7fffffffef164
(gdb) <enter> # repeat last action
(gdb) ... # keep pressing <enter>
(gdb) q
```

	+3	+2	+1	+0
0x0	xx	xx	xx	xx
0x4	xx	xx	xx	xx
...	xx	xx	xx	xx
0x7F...FE164	0xaa	0xbb	xx	xx
0x7F...FE168	0x12	0x34	0x56	0x78
0x7F...FE16C	'y'	'b'	xx	xx
0x7F...FE170	'e'	'h'	'\0'	'e'
0x7F...FE174	'\0'	'o'	'l'	'l'
0x7F...FE178	xx	xx	xx	xx
...	xx	xx	xx	xx
0xFFFFFFFFC	xx	xx	xx	xx





Structures, Revisited, Again

Read for HW

A "struct" is really an instruction to C on how to arrange a bunch of bytes in a bucket.

- Structs provide enough space for the data.
- C compilers often **align** the data with **padding**.

```
struct foo {  
    int32_t a;  
    char b;  
    struct foo *c;  
}
```

- For this struct, the actual layout on a 32b architecture would be as follows.
 - Note the 3 bytes of padding
 - `sizeof(struct foo) == 12`

+3	+2	+1	+0
	4 bytes for a		
unused	unused	unused	1 byte for b
	4 bytes for c		